

21.08.2019

CONFIDENTIAL

MODERN DATA WAREHOUSE

Microsoft Azure



MISTER SPEX

SO MANY DATA SOURCES



Google
BigQuery



{ j s o n }

REST





EVANDRO MUCHINSKI

- Data Engineer at Mister Spex
- 7 years as Microsoft Certified Trainer
- MCSE, MCSA...
- Brazilian living in Berlin for almost 3 years
- Working with data driven projects since 2006



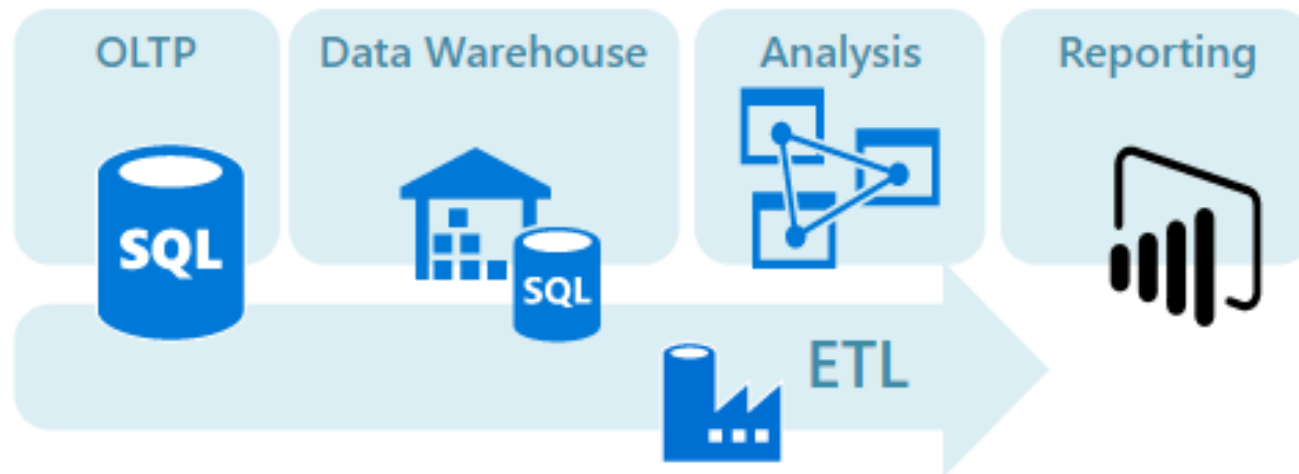
LinkedIn <https://www.linkedin.com/in/evandro-muchinski/>

AGENDA

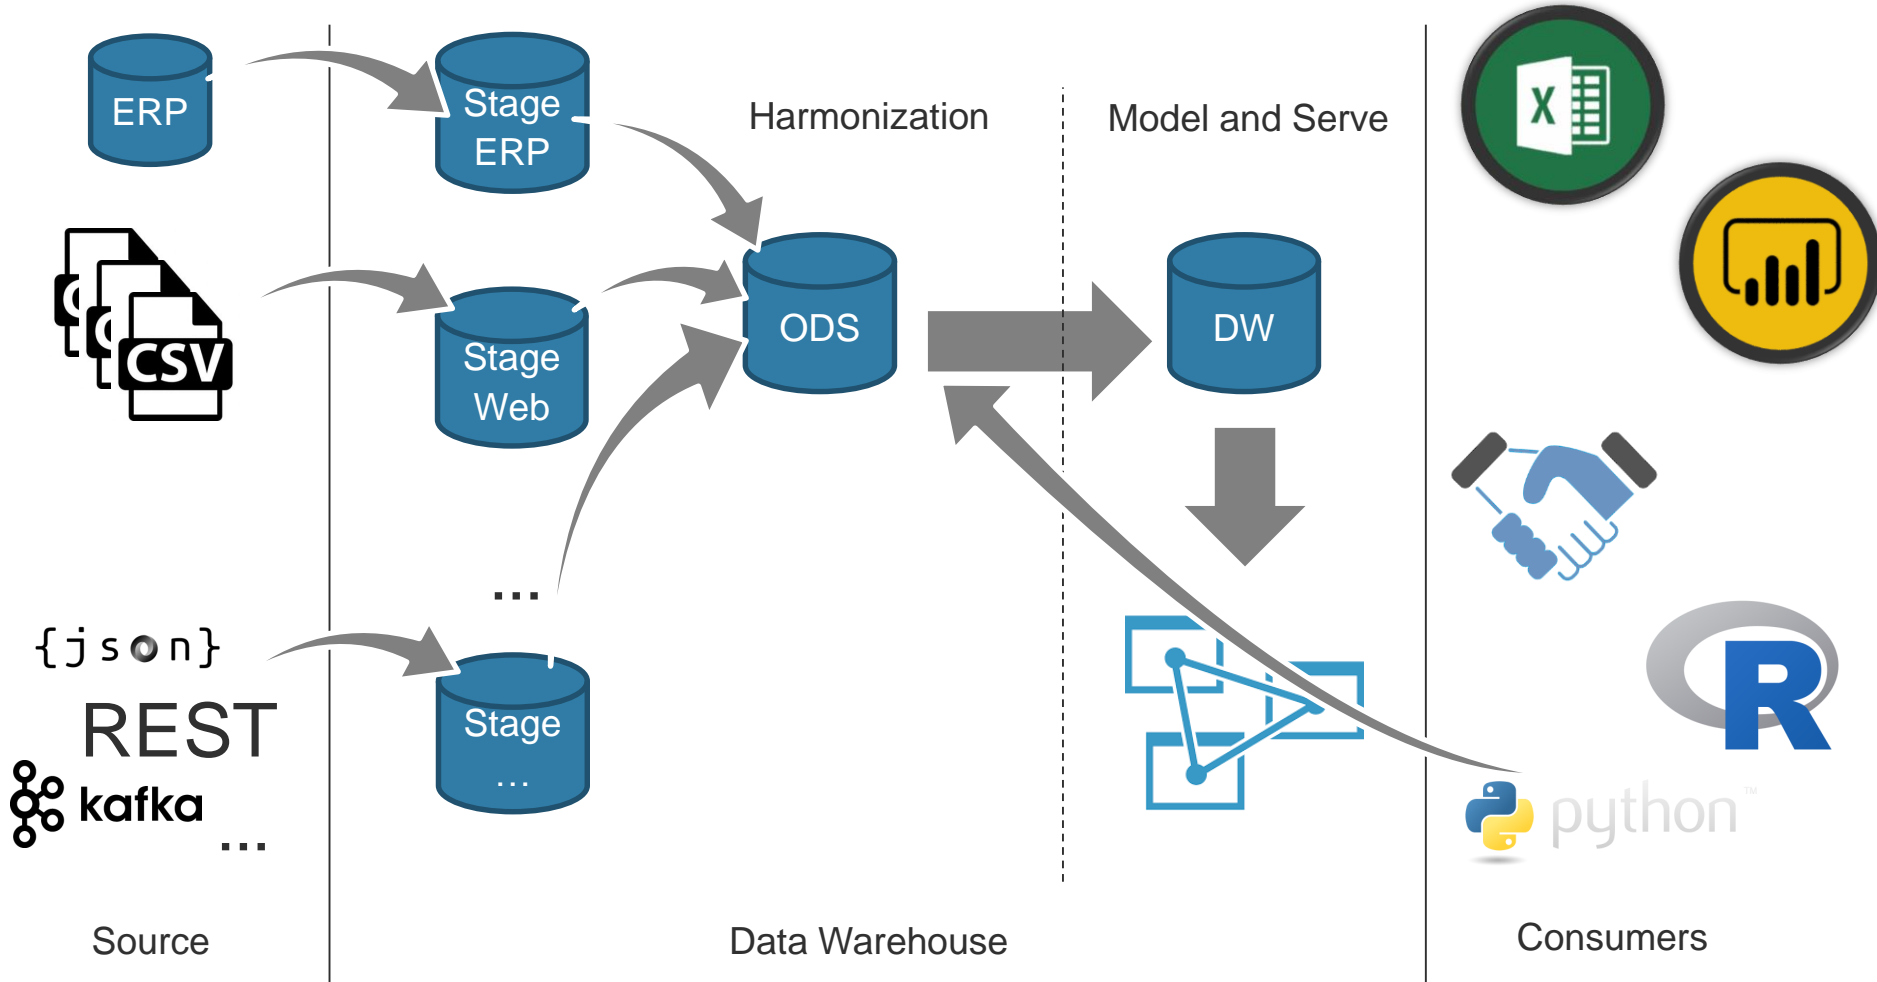
1. PROBLEM
2. TRADITIONAL DATA WAREHOUSE ARCHITECTURE
3. MODERN DATA WAREHOUSE
4. TOOLSET OVERVIEW
5. REFERENCES



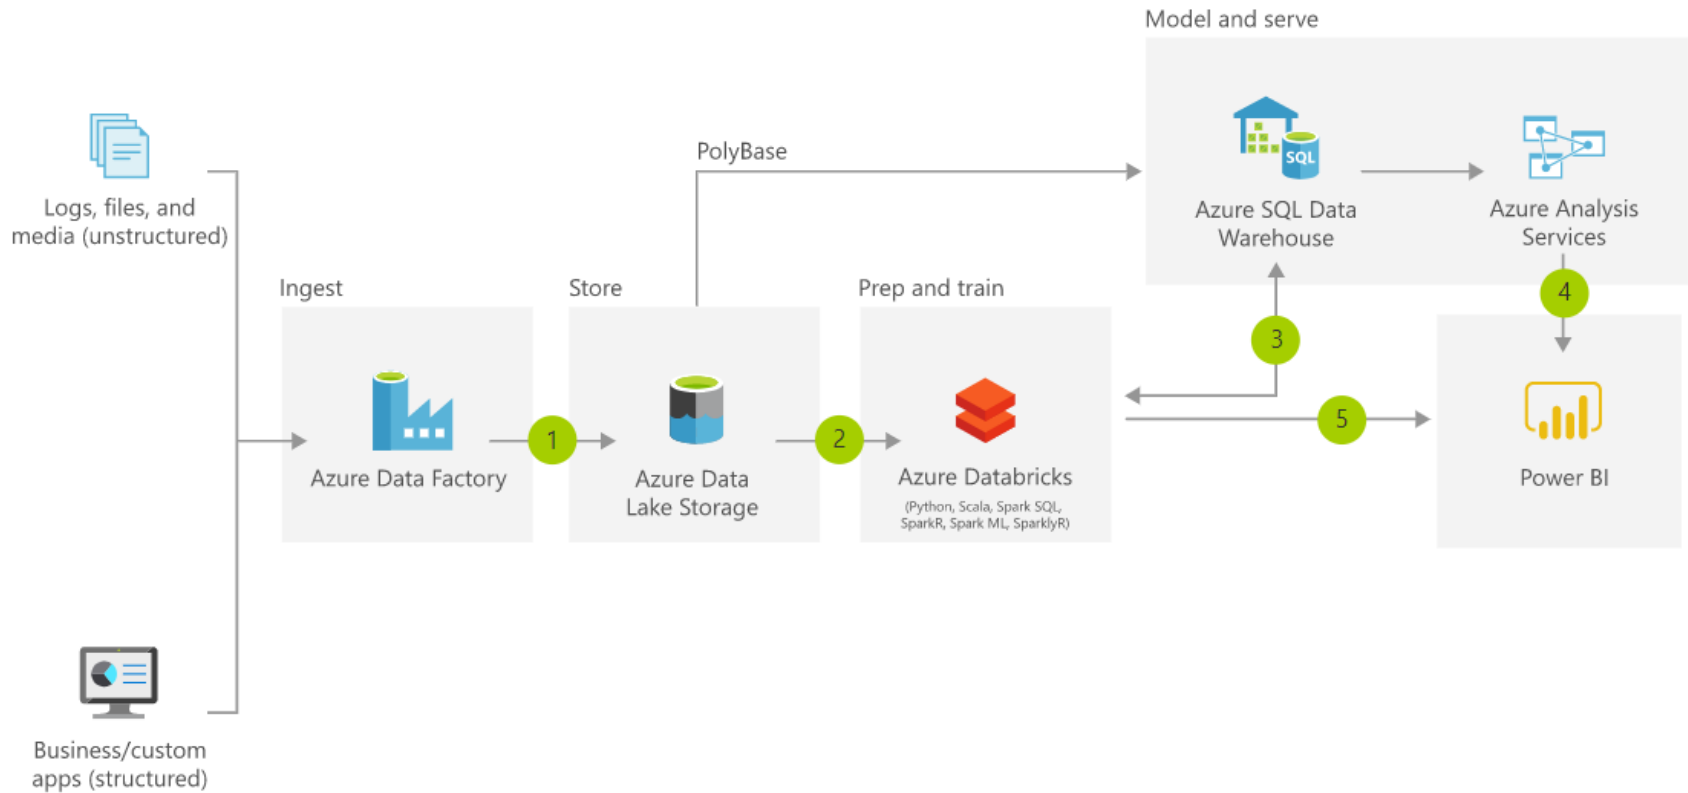
GENERAL ARCHITECTURE



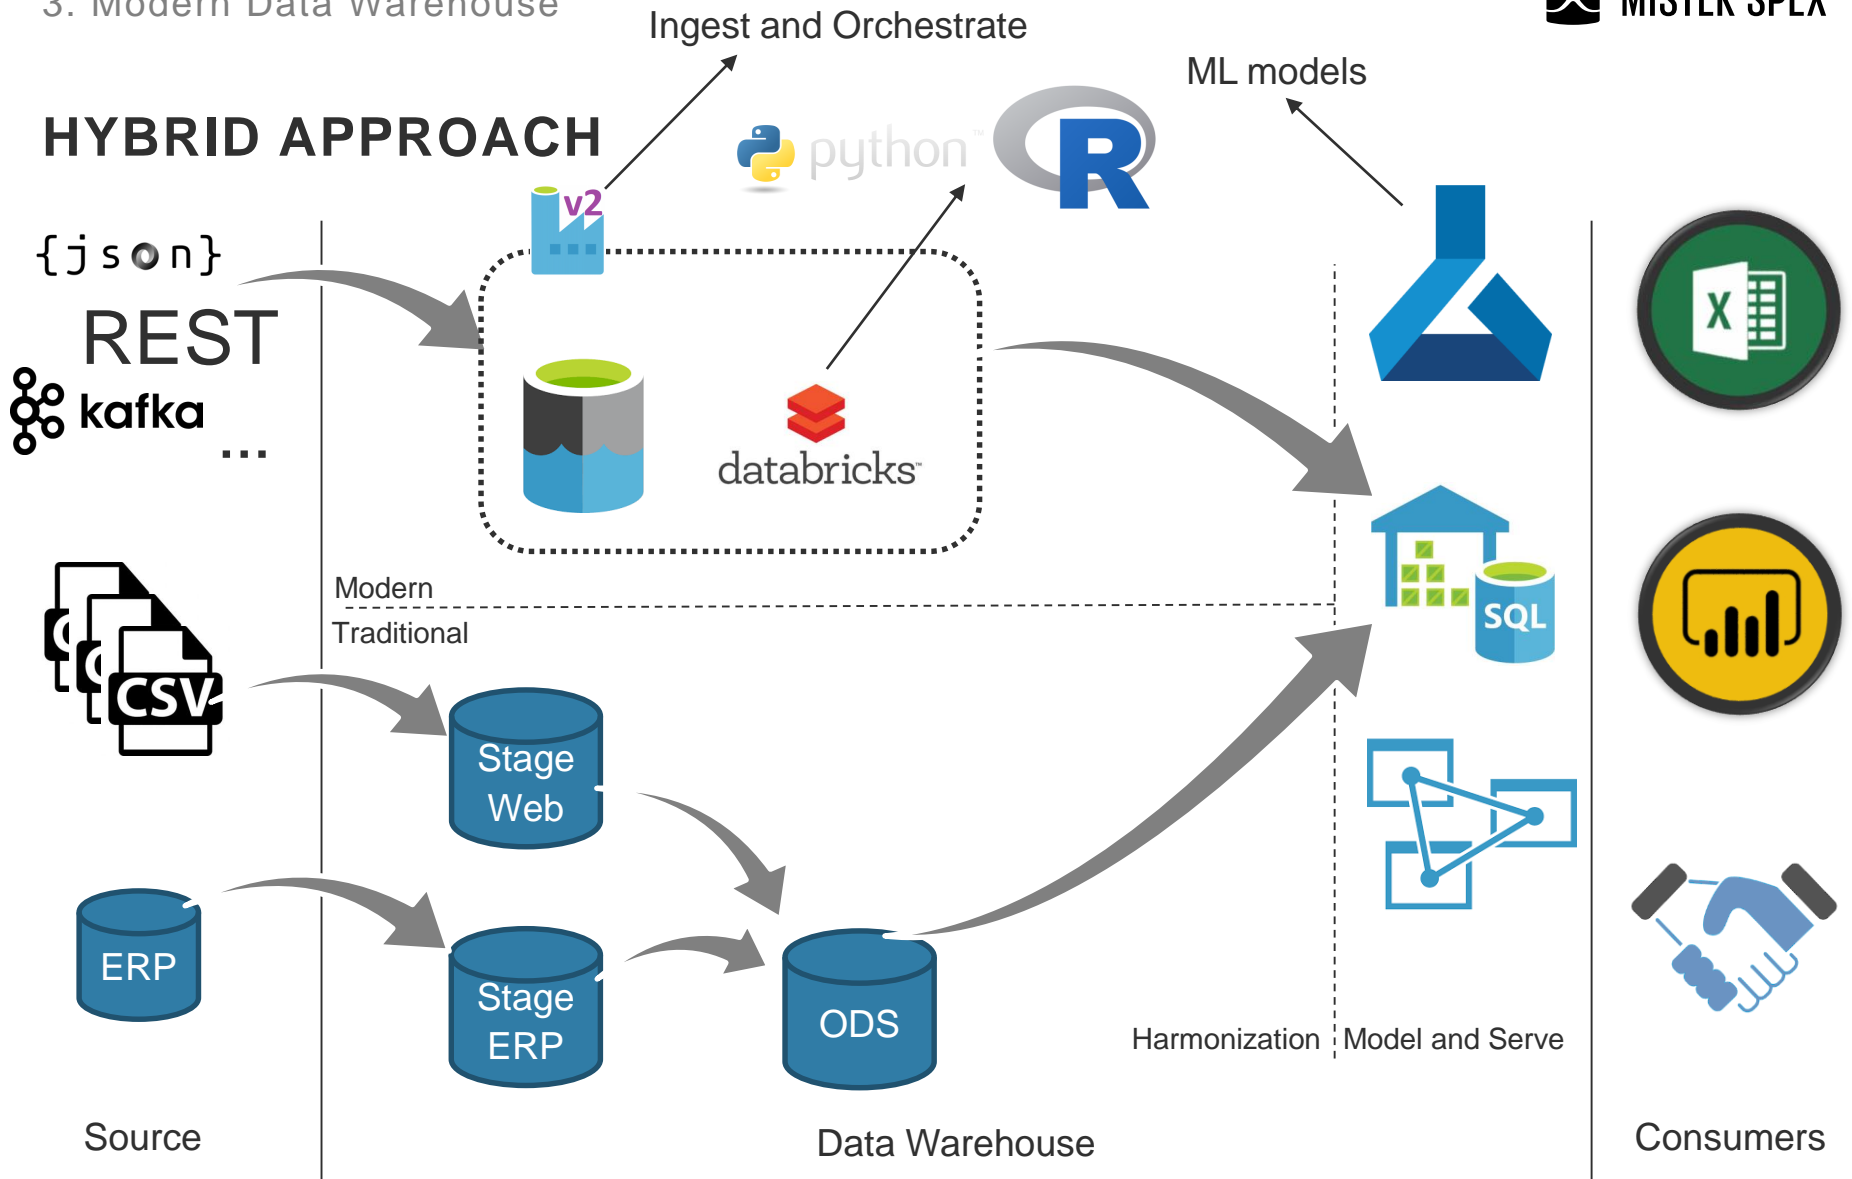
SOMETHING CLOSE TO WHAT WE HAVE



GENERAL ARCHITECTURE



3. Modern Data Warehouse

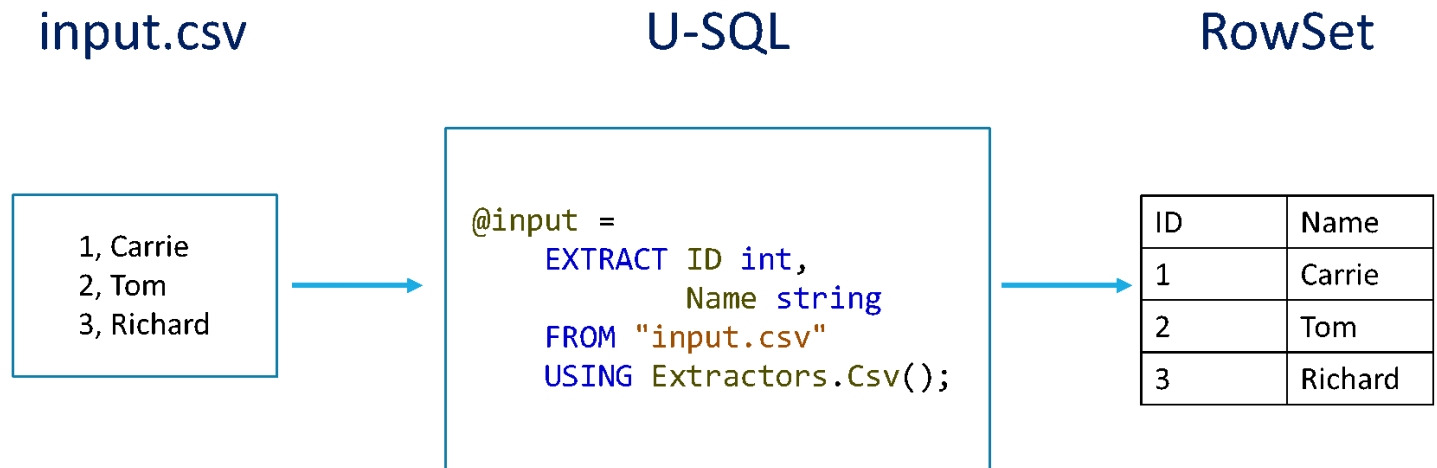


NEW TOOLS AVAILABLE ON AZURE

- **Data Lake**
- **Data Factory V2**
- **Data Bricks**
- **Data Warehouse**
- **Machine Learning Services**

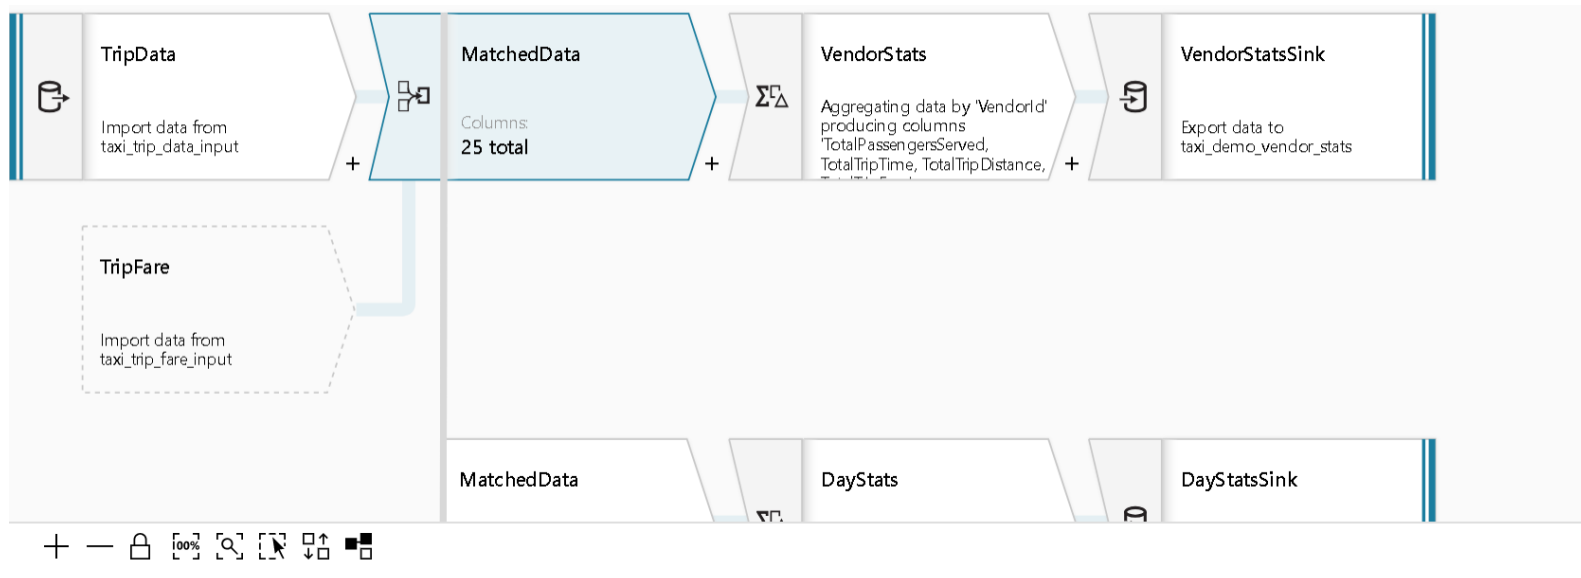
DATA LAKE

- High scalable storage
- Hadoop compatible access
- Files managed in directories and subdirectories
- Permission can be granted to directories or files
- U-SQL compatible



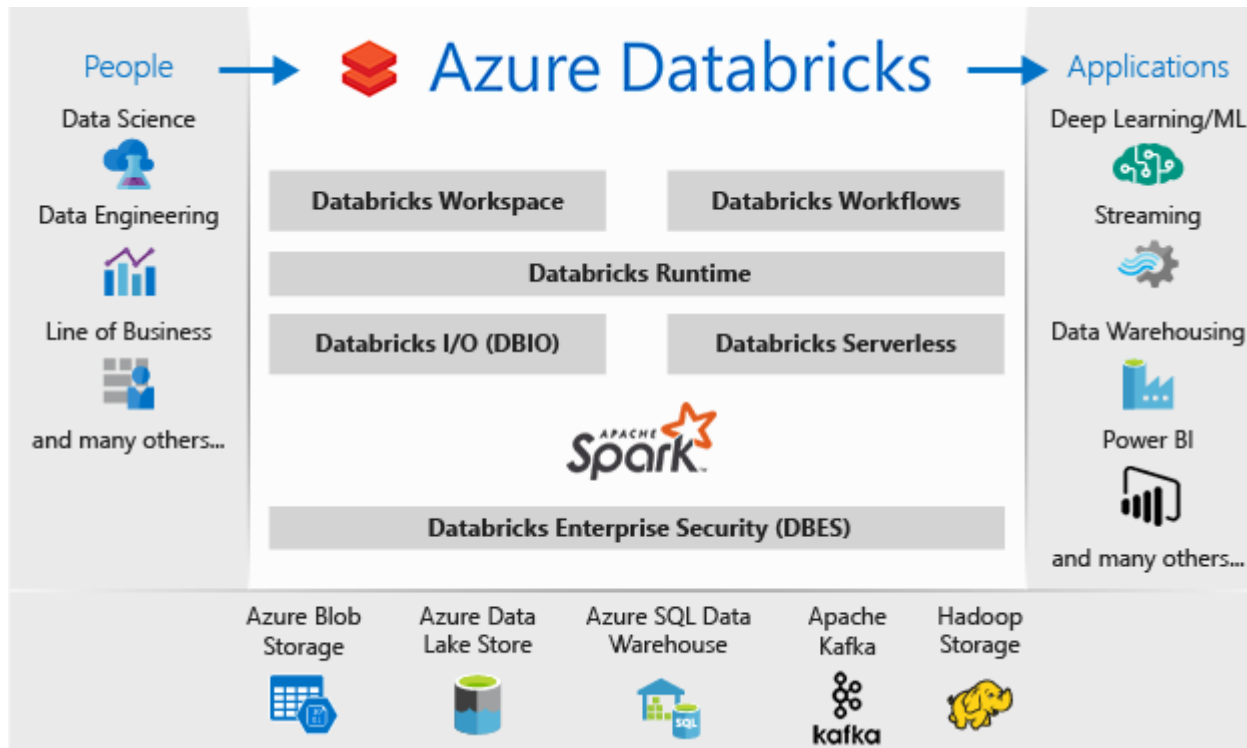
DATA FACTORY V2

- Ingest and Orchestrate tool
- 80+ connectors
- Integrated with Git
- Develop with Code or Interface
- Compatible with Azure Data Bricks
- Dataflows = Spark Executables



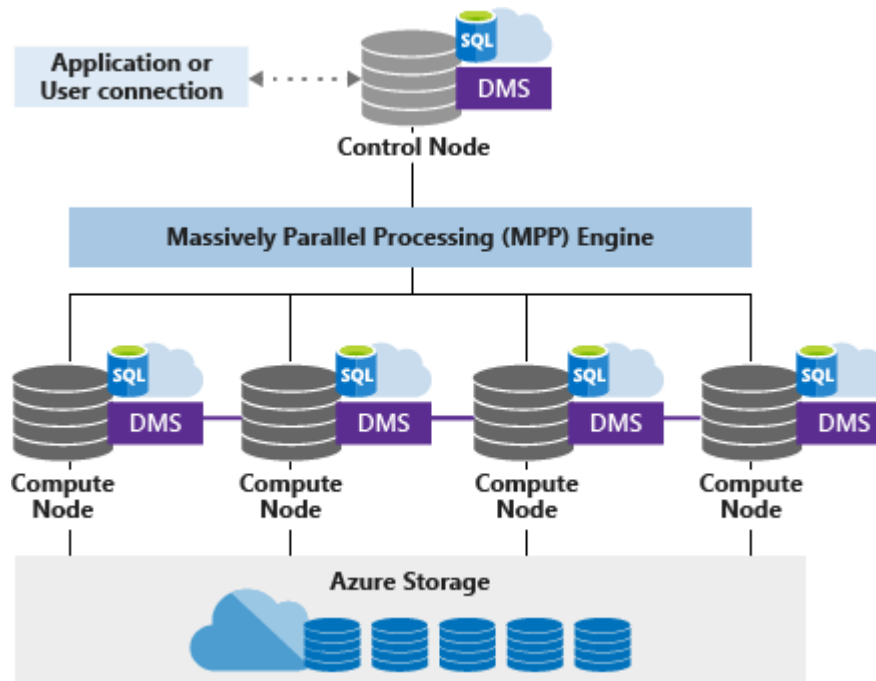
DATA BRICKS

- Runs Python, R, Scala, SQL



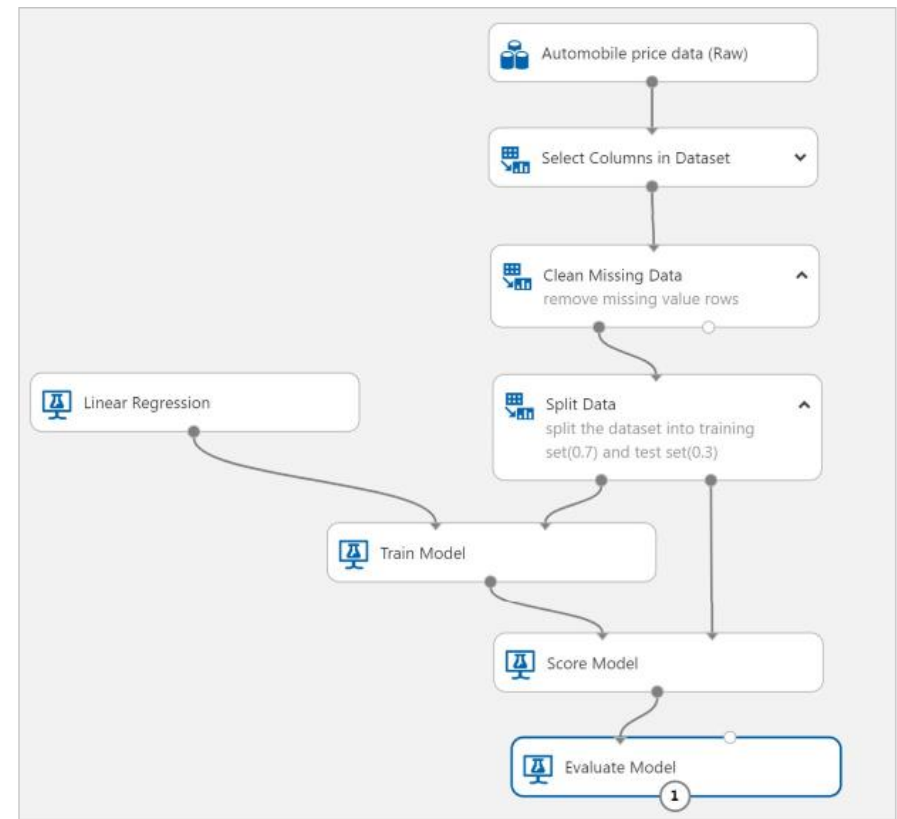
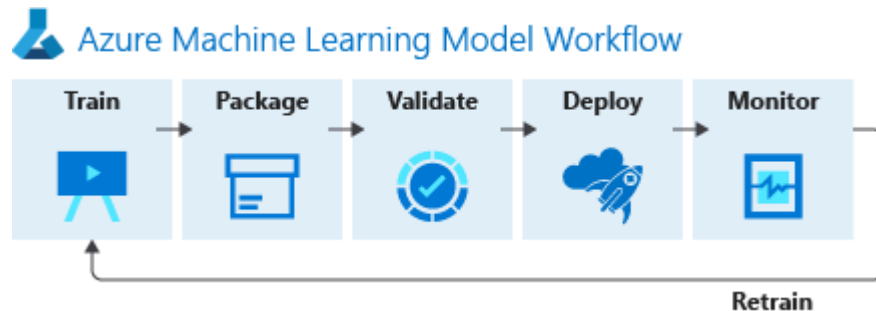
DATA WAREHOUSE

- Massive Parallel Processing column-store database cluster



MACHINE LEARNING SERVICES

- Automate and scale Data Science projects
- Work with Python only
- More on the next session 😊



LEARN MORE AT:

[Azure Data Architecture Guide](#)

[Microsoft Build 2019](#)

[Learning path for Azure Data Engineer](#)

THANK YOU!

